

DL Dev Course: Week 03b

What are TF Estimators ?

- A high level abstraction API for writing TensorFlow models
- What Google wants to become the standard for writing most TF models
- They consist of a number of key parts and come in both premade and custom versions
- Take inspiration from SKLearn and from Keras

Flexible: High level APIs

Low-Level Python API

```
import numpy as np
import tensorflow as tf

W = tf.Variable([.3], tf.float32)
b = tf.Variable([-3], tf.float32)
x = tf.placeholder(tf.float32)
linear_model = W * x + b
y = tf.placeholder(tf.float32)
loss = tf.reduce_sum(tf.square(linear_model - y))
optimizer = tf.train.GradientDescentOptimizer(0.01)
train = optimizer.minimize(loss)
x_train = [1,2,3,4]
y_train = [0,-1,-2,-3]
init = tf.global_variables_initializer()
sess = tf.Session()
sess.run(init) # reset values to wrong
for i in range(1000):
    sess.run(train, {x:x_train, y:y_train})
curr_W, curr_b, curr_loss = sess.run([W, b, loss], {x:x_train, y:y_train})
```

High-Level Python API

```
import tensorflow as tf
import numpy as np

features = [tf.contrib.layers.real_valued_column("x", dimension=1)]
estimator = tf.contrib.learn.LinearRegressor(feature_columns=features)
x = np.array([1., 2., 3., 4.])
y = np.array([0., -1., -2., -3.])
input_fn = tf.contrib.learn.io.numpy_input_fn({"x":x}, y, batch_size=4,
                                              num_epochs=1000)

estimator.fit(input_fn=input_fn, steps=1000)
estimator.evaluate(input_fn=input_fn)
```

Flexible: High level APIs

Low-Level Python API

```
import numpy as np
import tensorflow as tf
```

15 lines

High-Level Python API

```
import tensorflow as tf
import numpy as np
```

```
features = [tf.contrib
estimator = tf.contrib
x = np.array([1., 2.,
y = np.array([0., -1.,
input_fn = tf.contrib.
```

```
estimator.fit(input_features, output)
estimator.evaluate(input_features, output)
```

7 lines

The Old Way 01

- feed_dict
- Train loop
- Low level ops
- Stuff all over the place
- Lots of lines of code

The Old Way 02

- Lots of Frameworks
- Slim
- TFLearn
- Sugar Tensor
- Etc. etc.

Python Frontend

C++

Java

Go

...

TensorFlow Distributed Execution Engine

CPU

GPU

Android

iOS

XLA

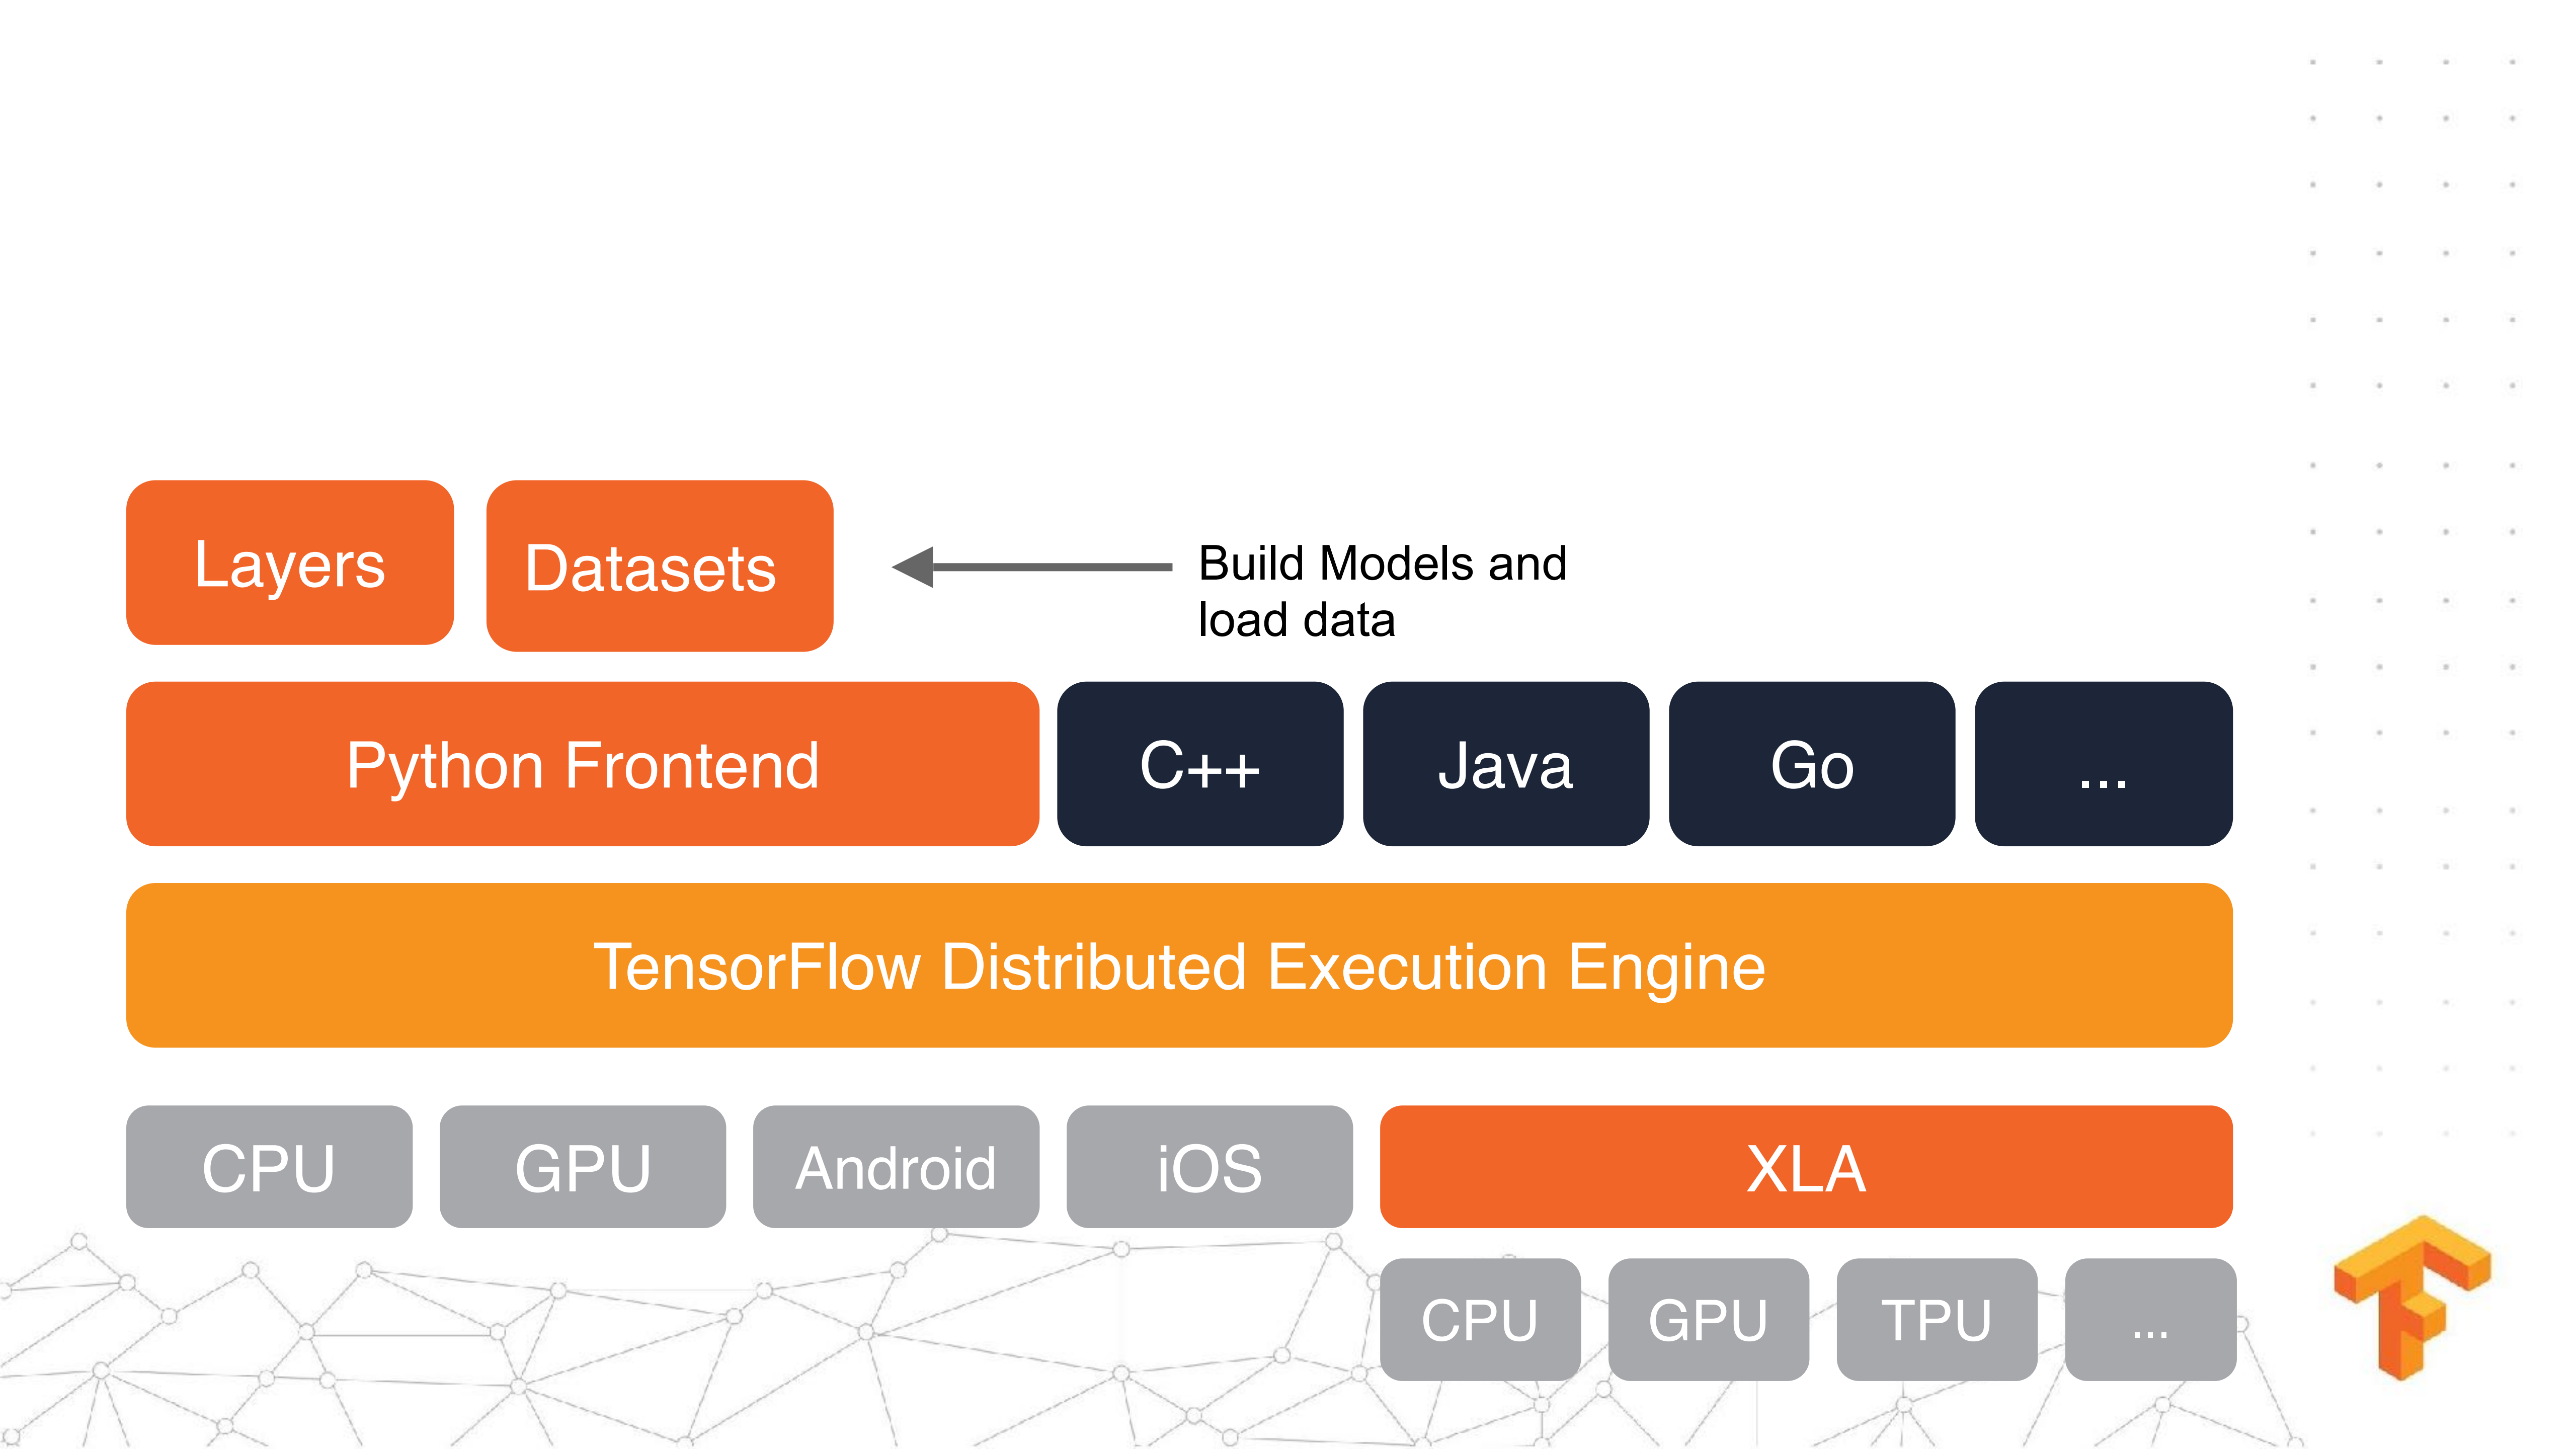
CPU

GPU

TPU

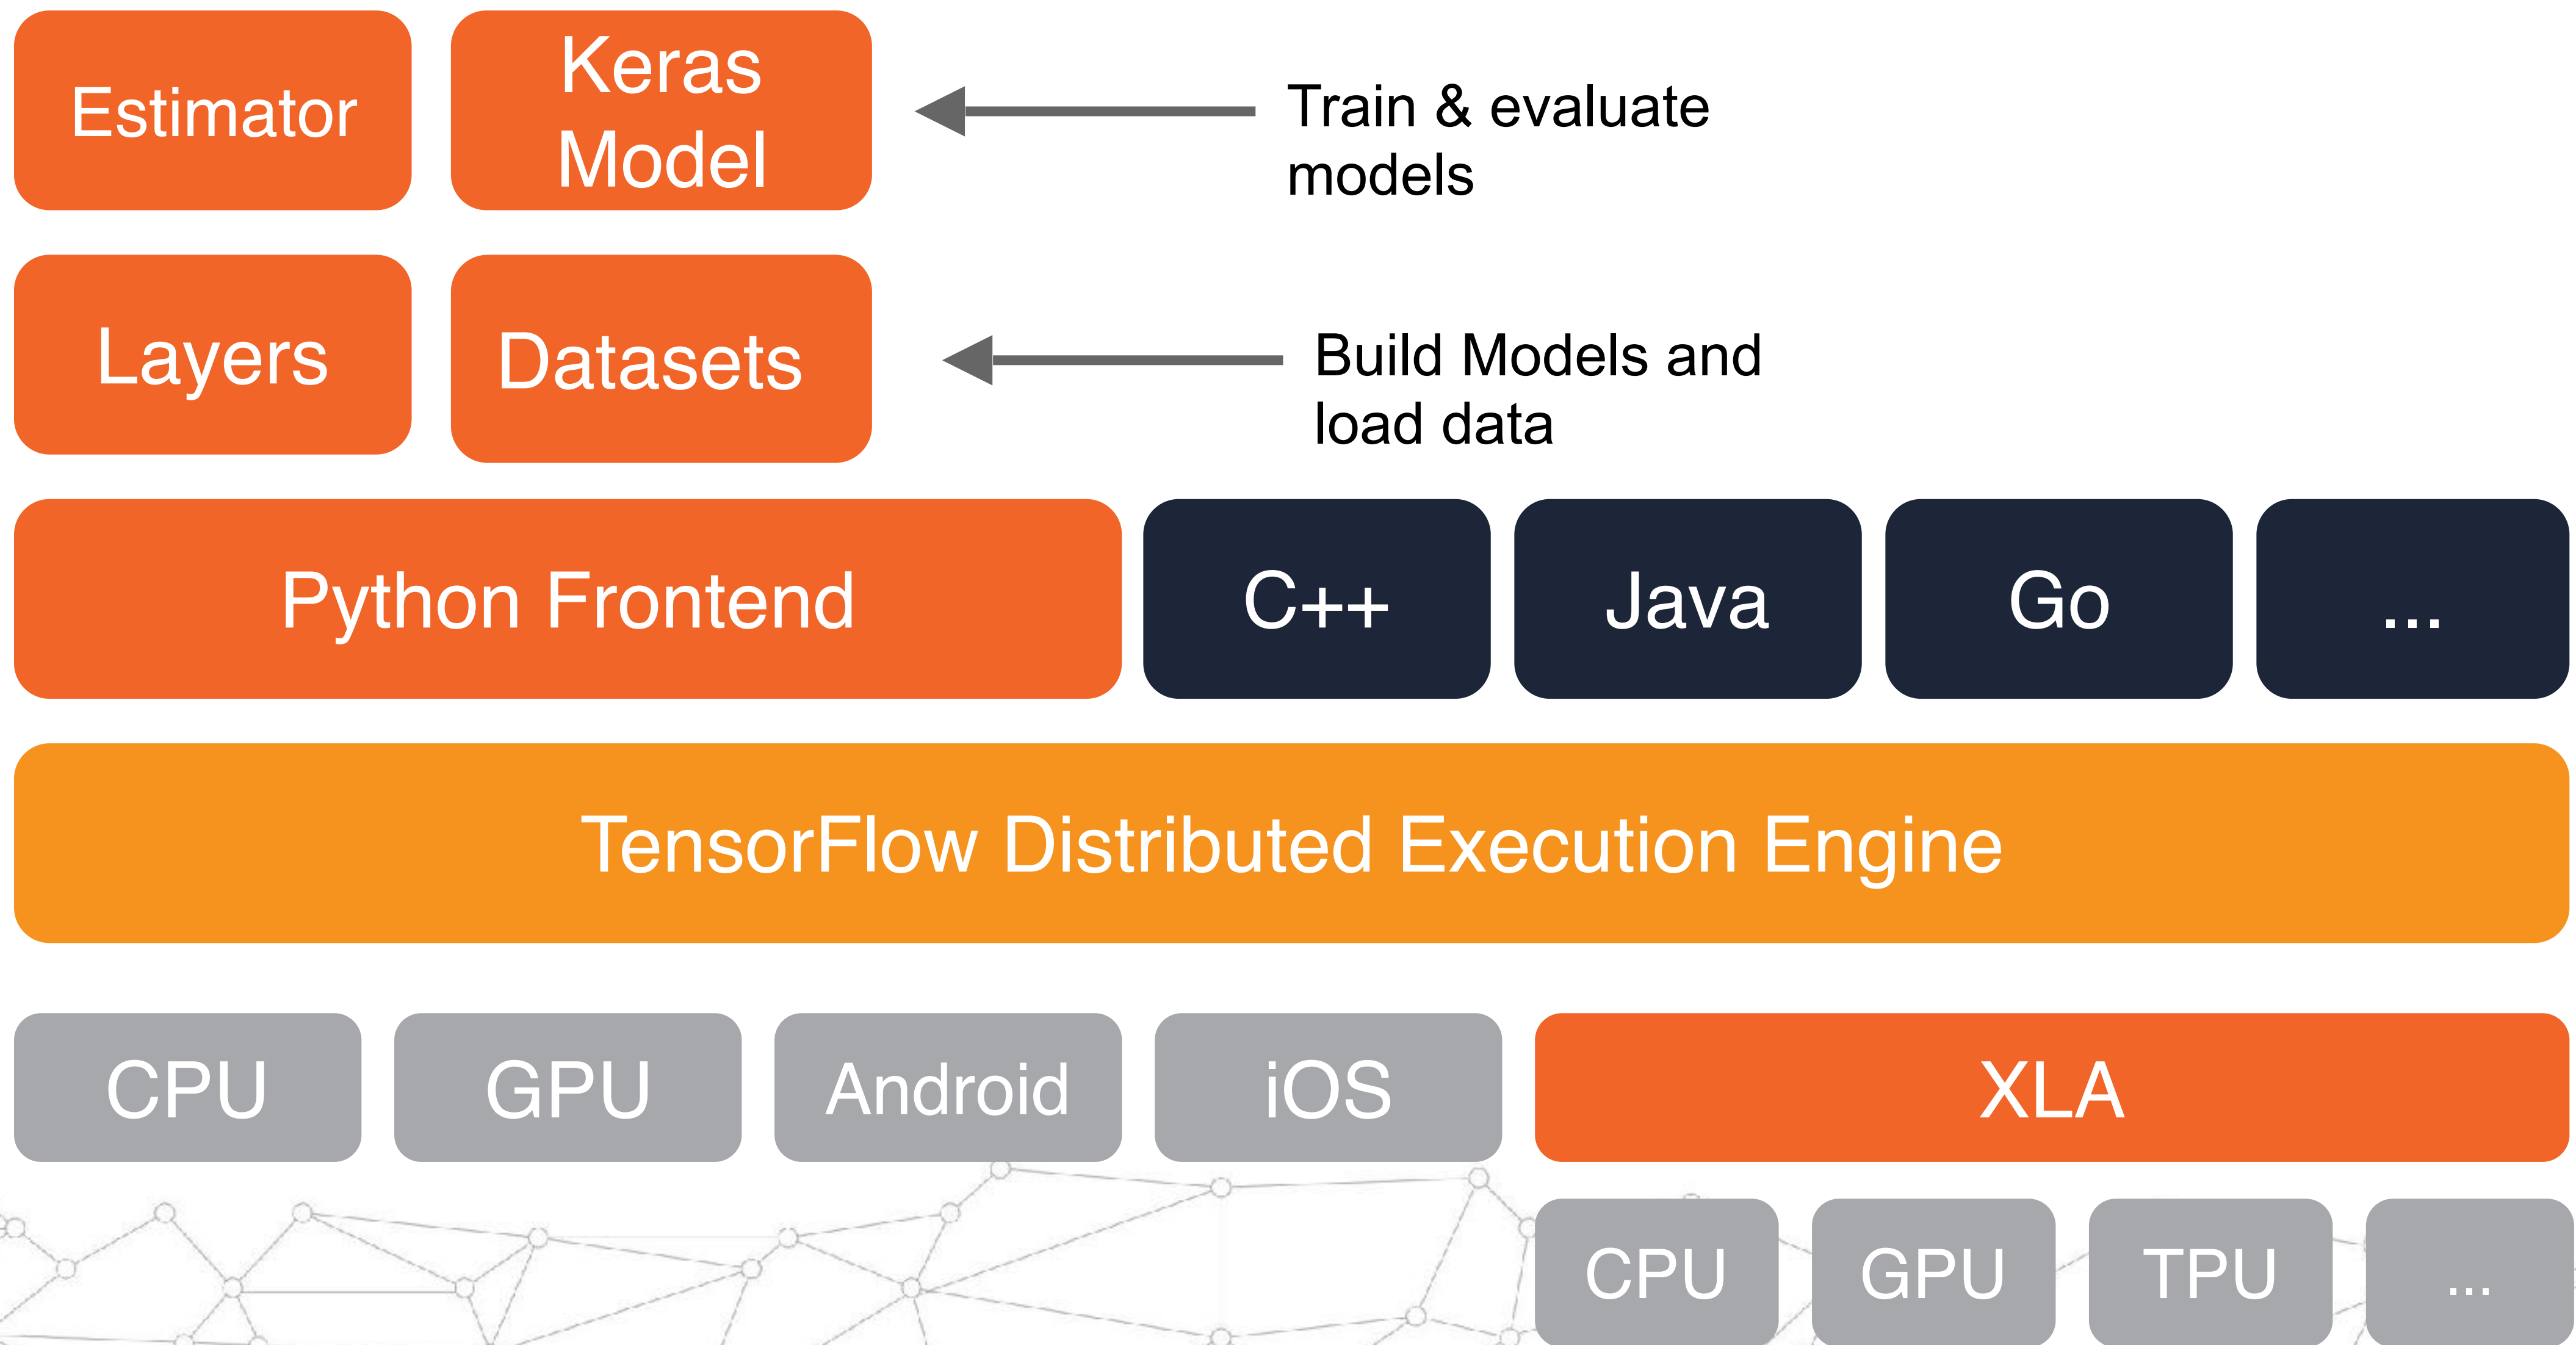
...

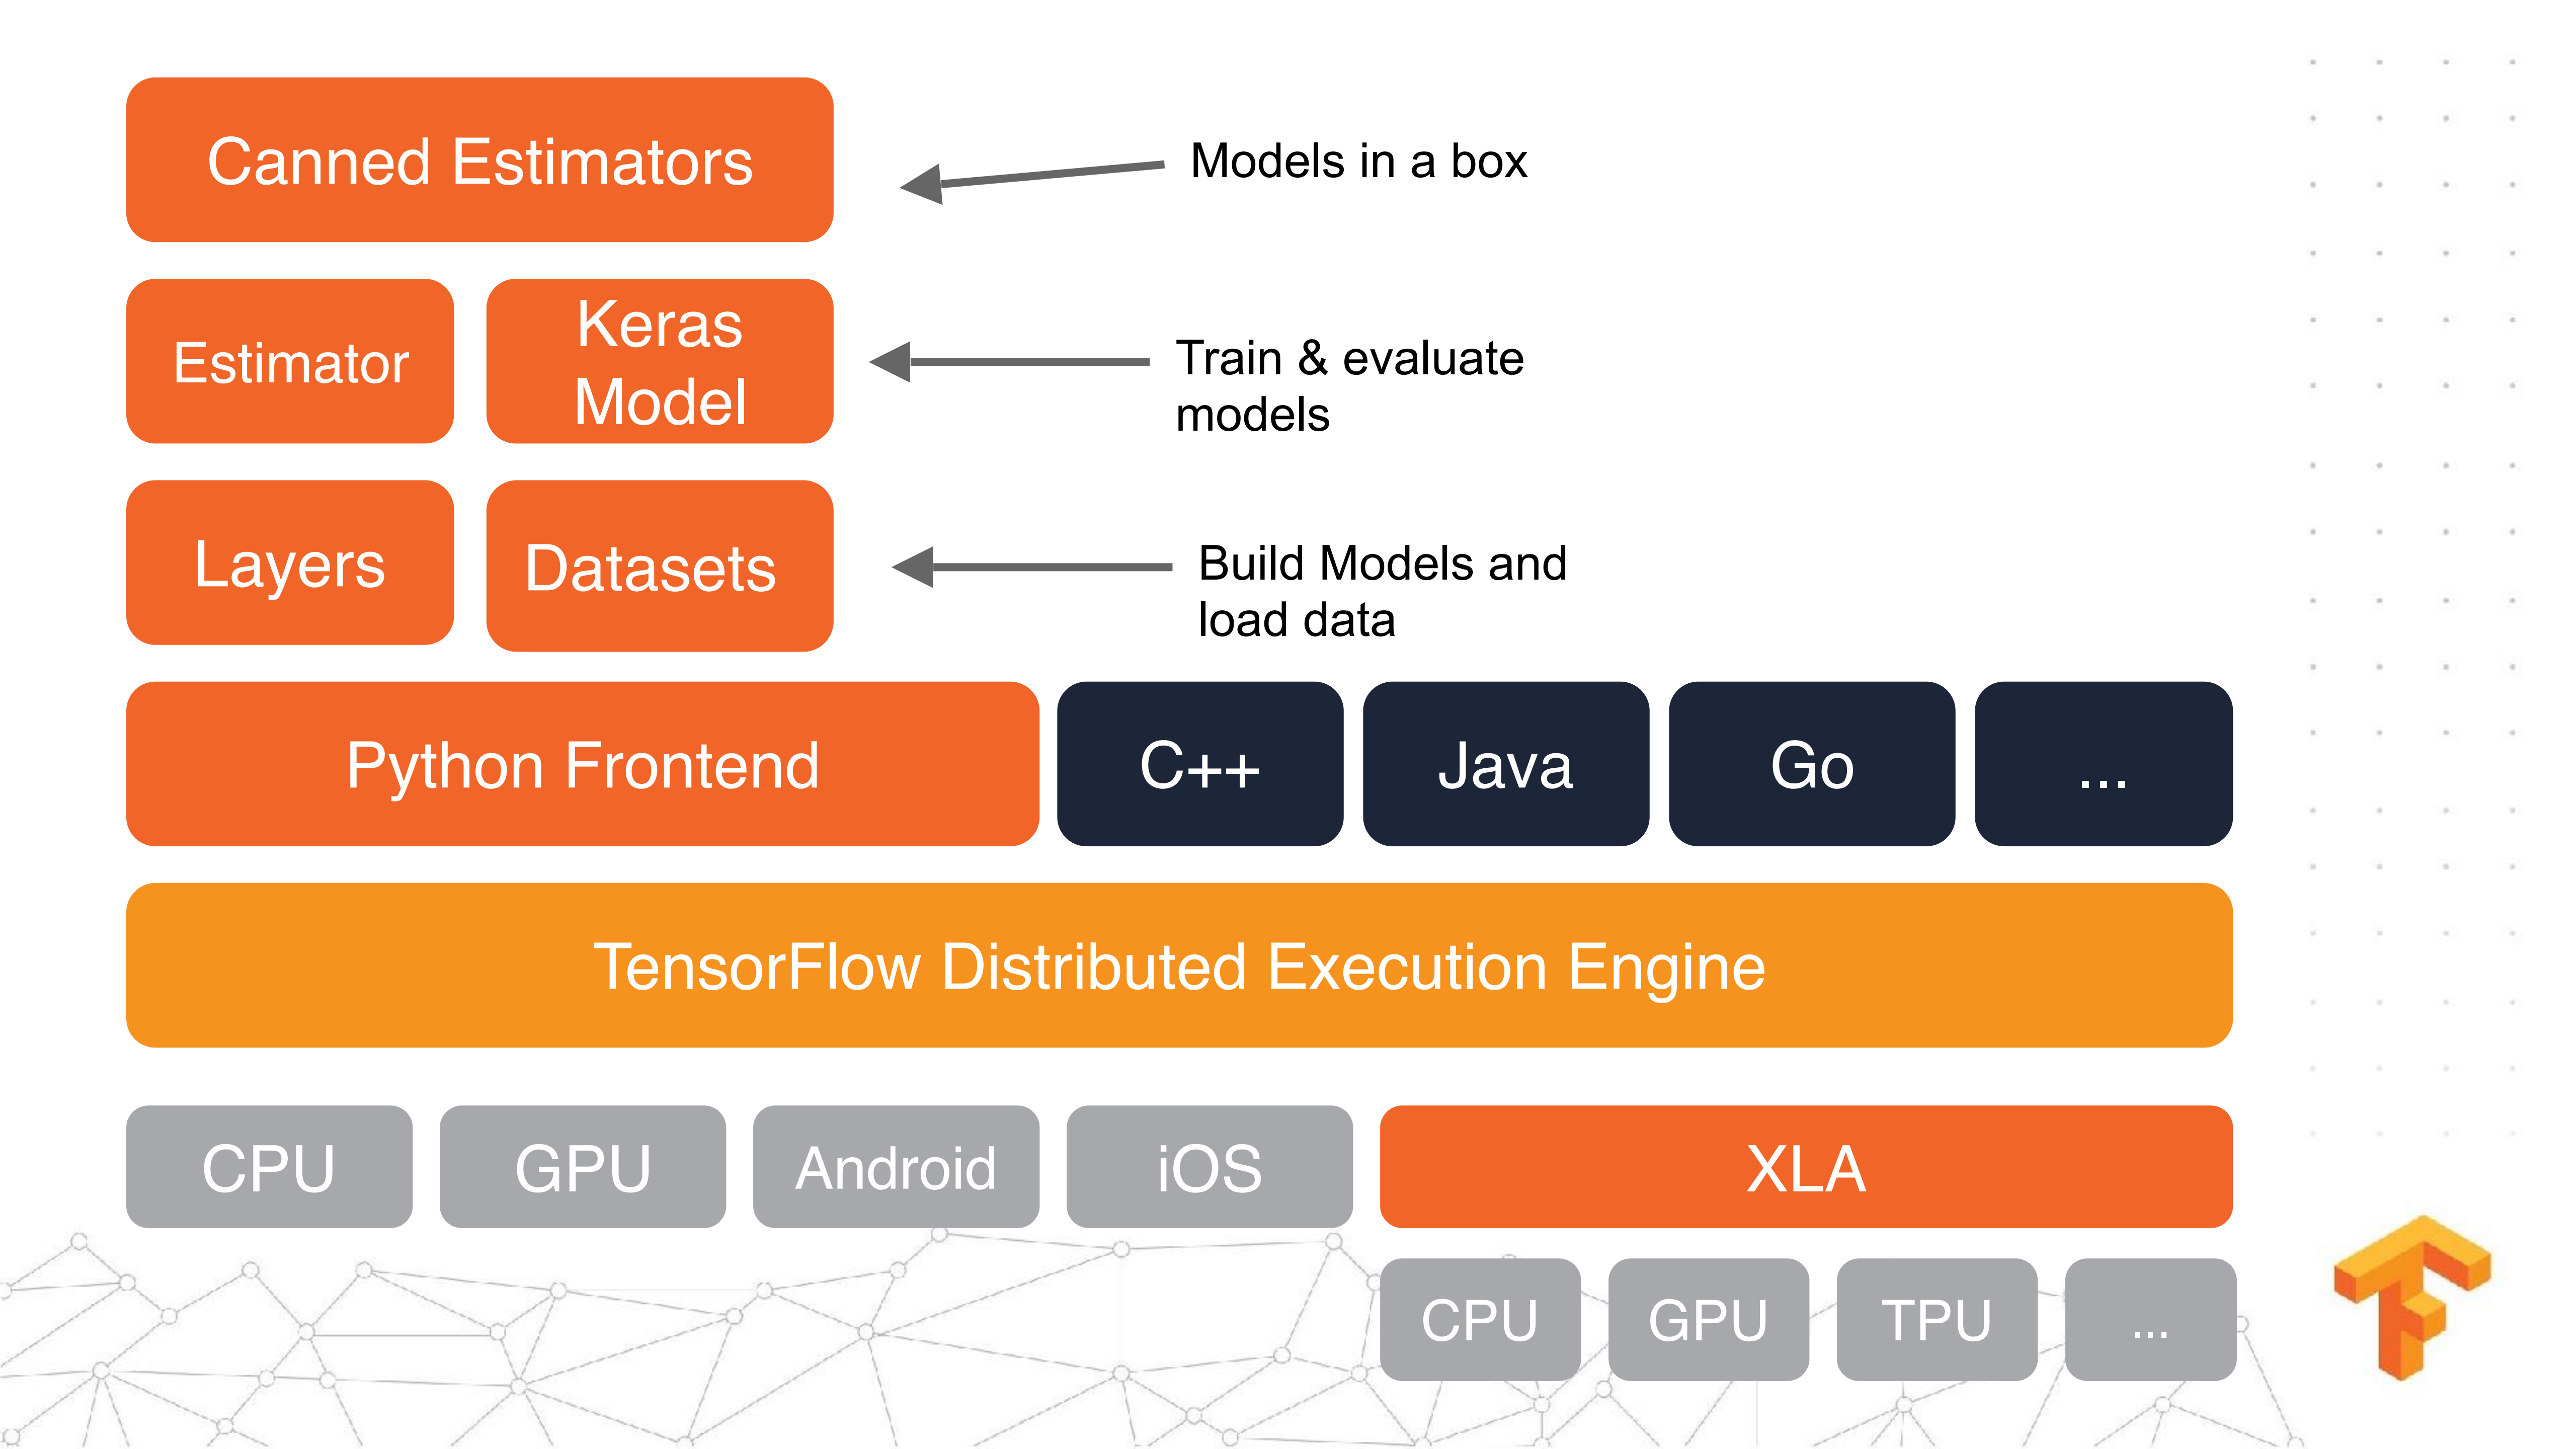




TF Layers


- Pre made layer that represent some of the more common types of NN operations
- `tf.layers.conv2d(inputs=image_batch, filters=32, kernel_size=[3, 3], padding='same', activation=tf.nn.relu)`
- `tf.layers.dense(inputs=tensor_shape, units=128, activation=tf.nn.relu)`
- They take inspiration from Keras

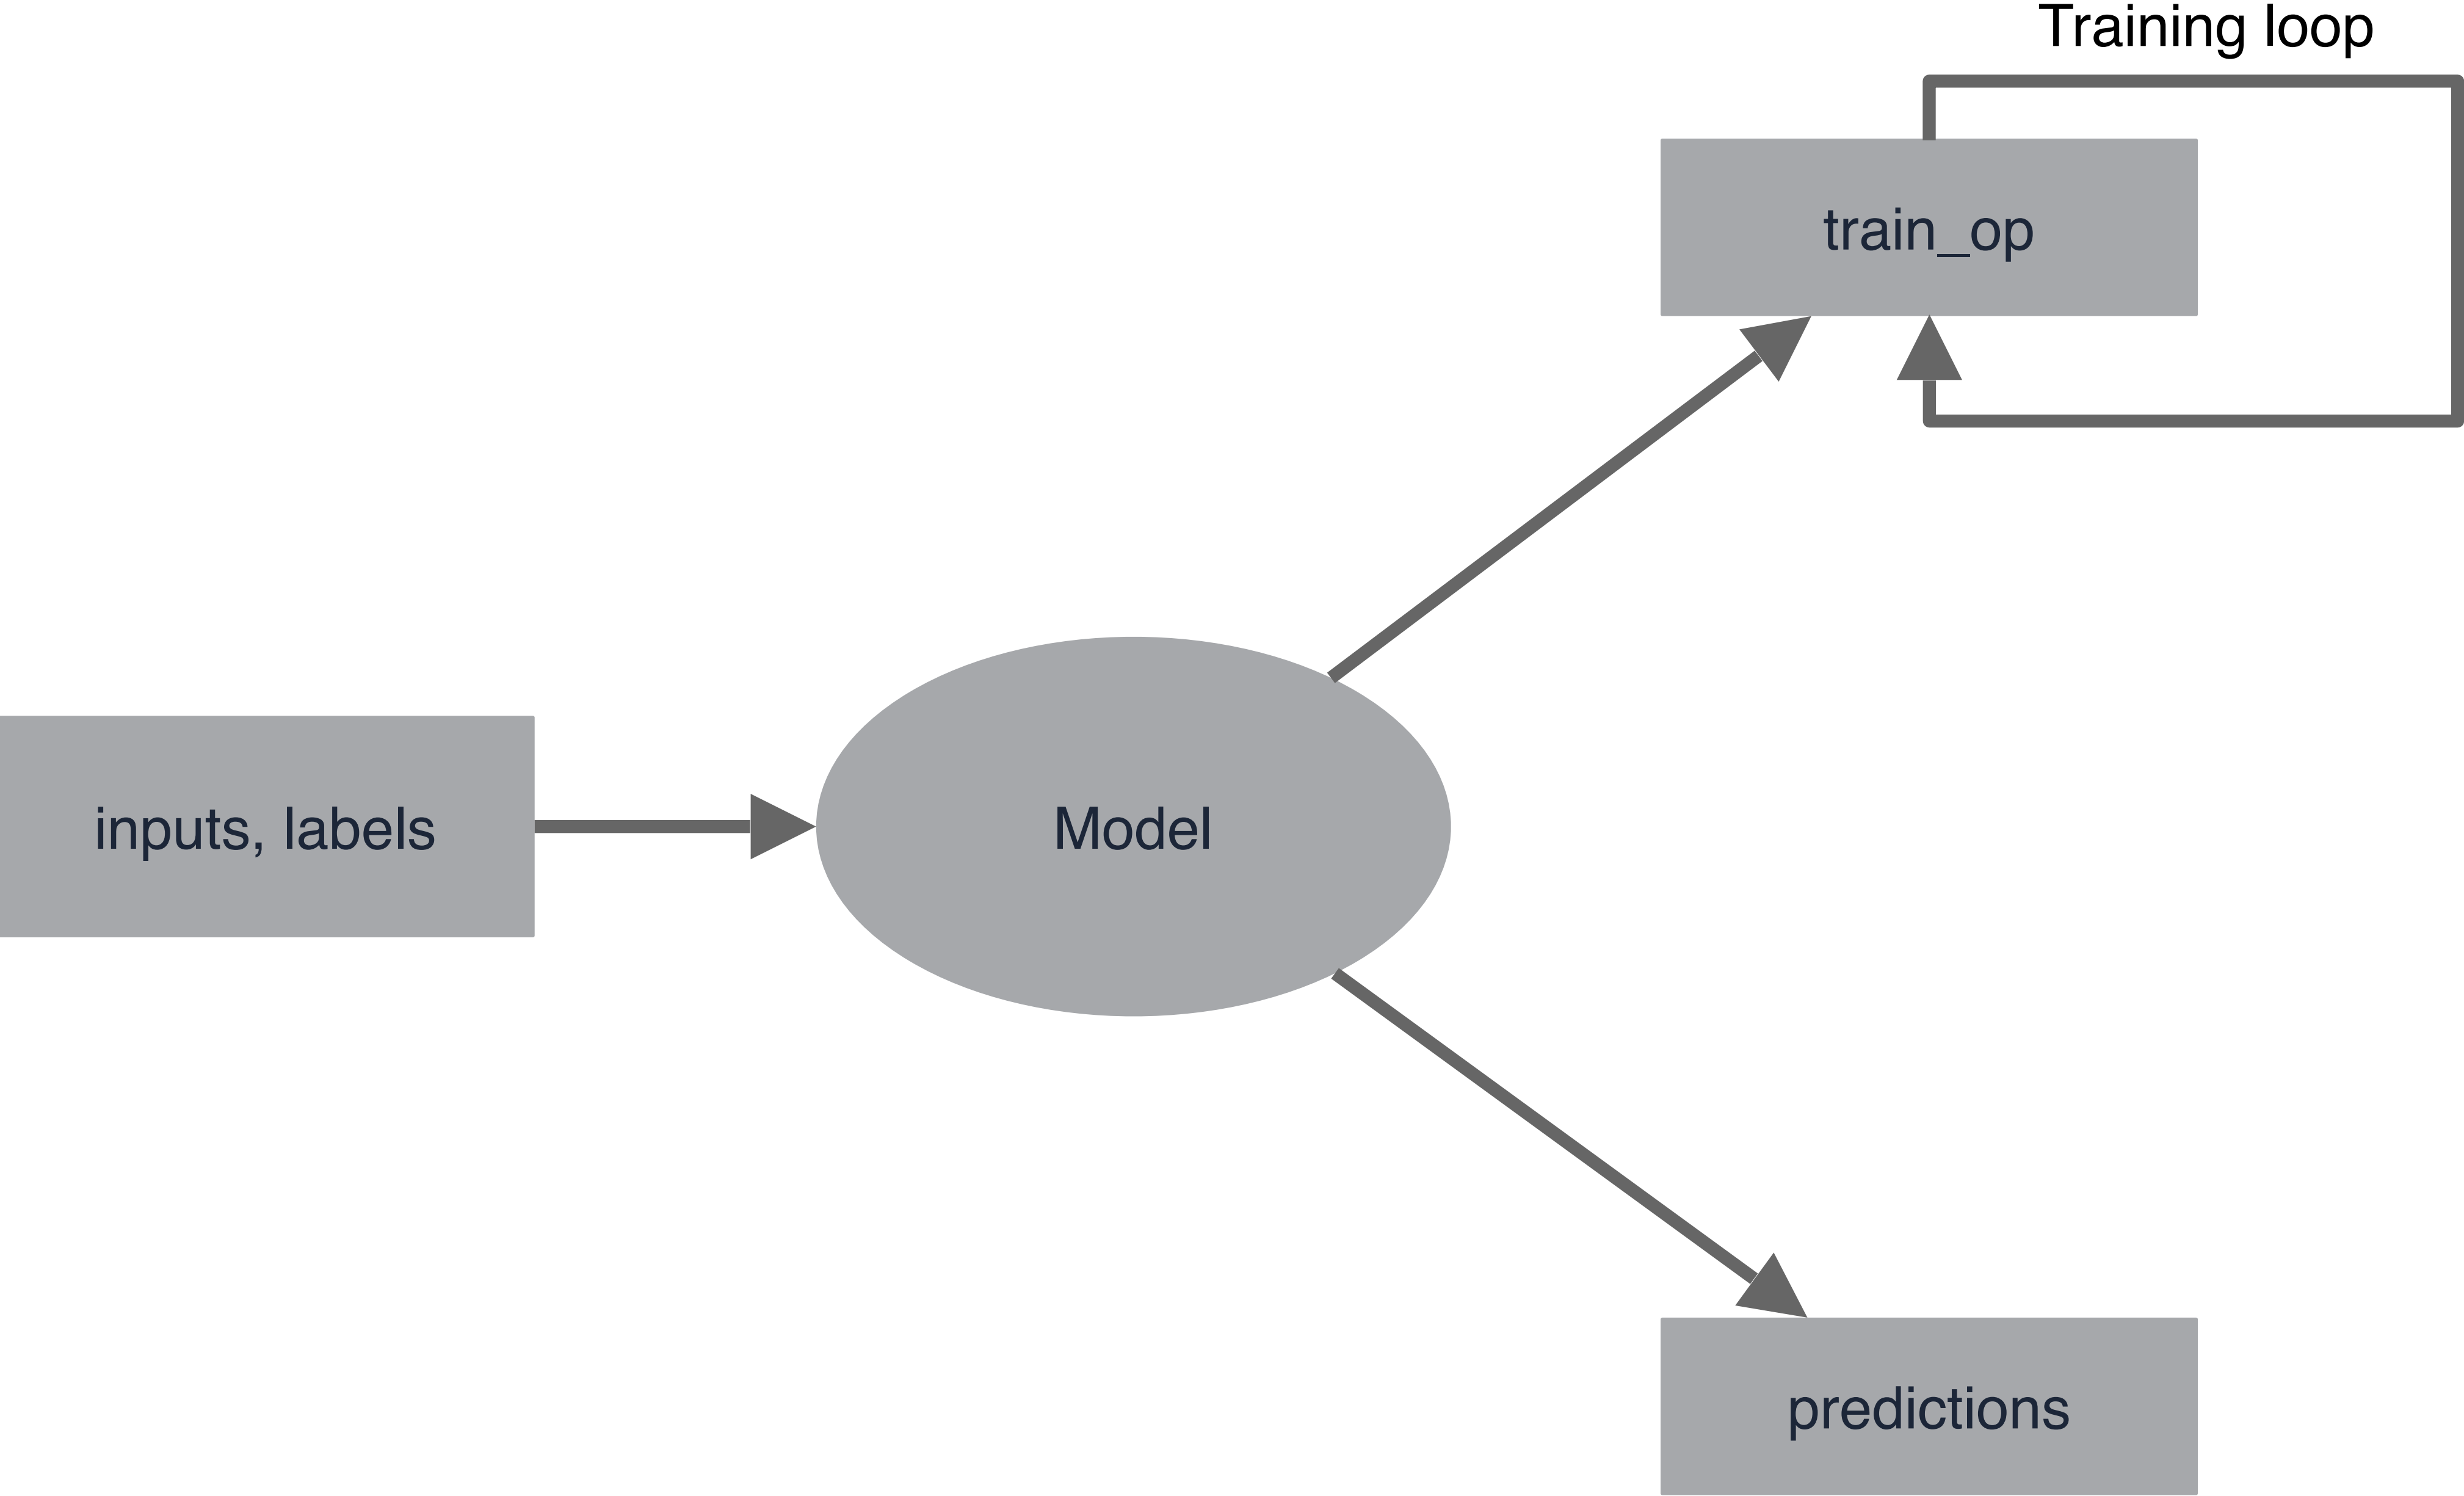


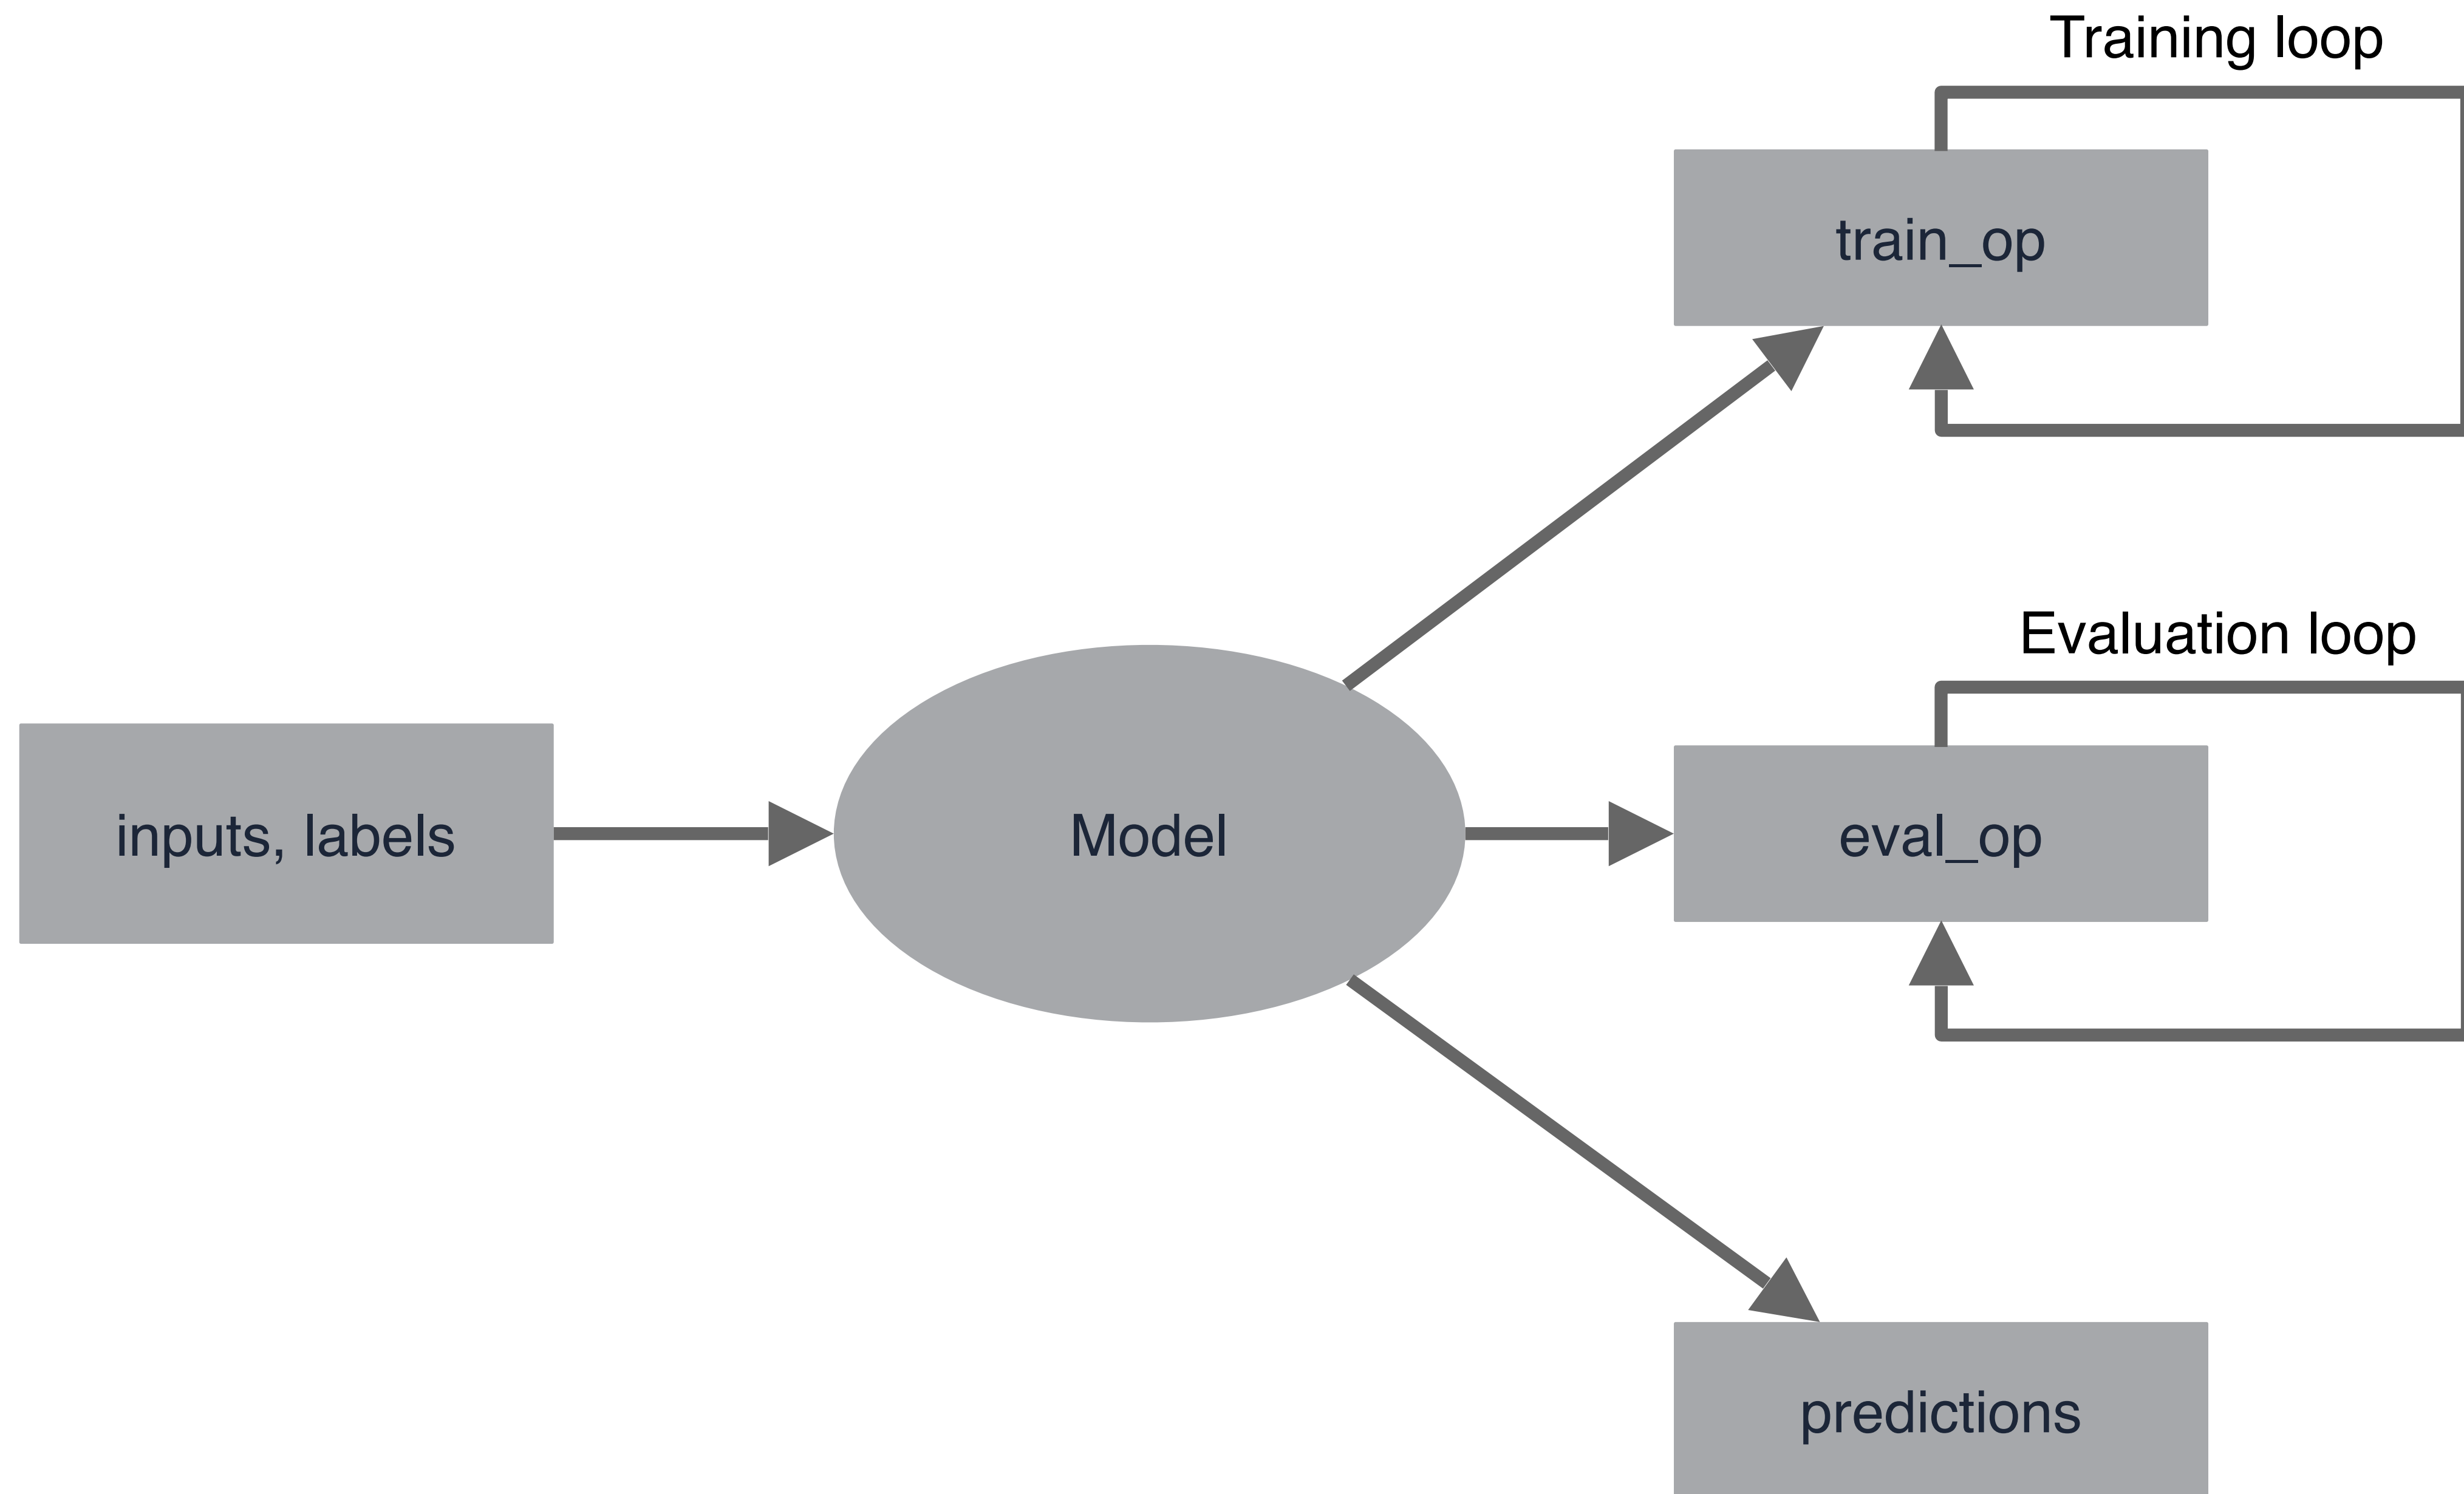


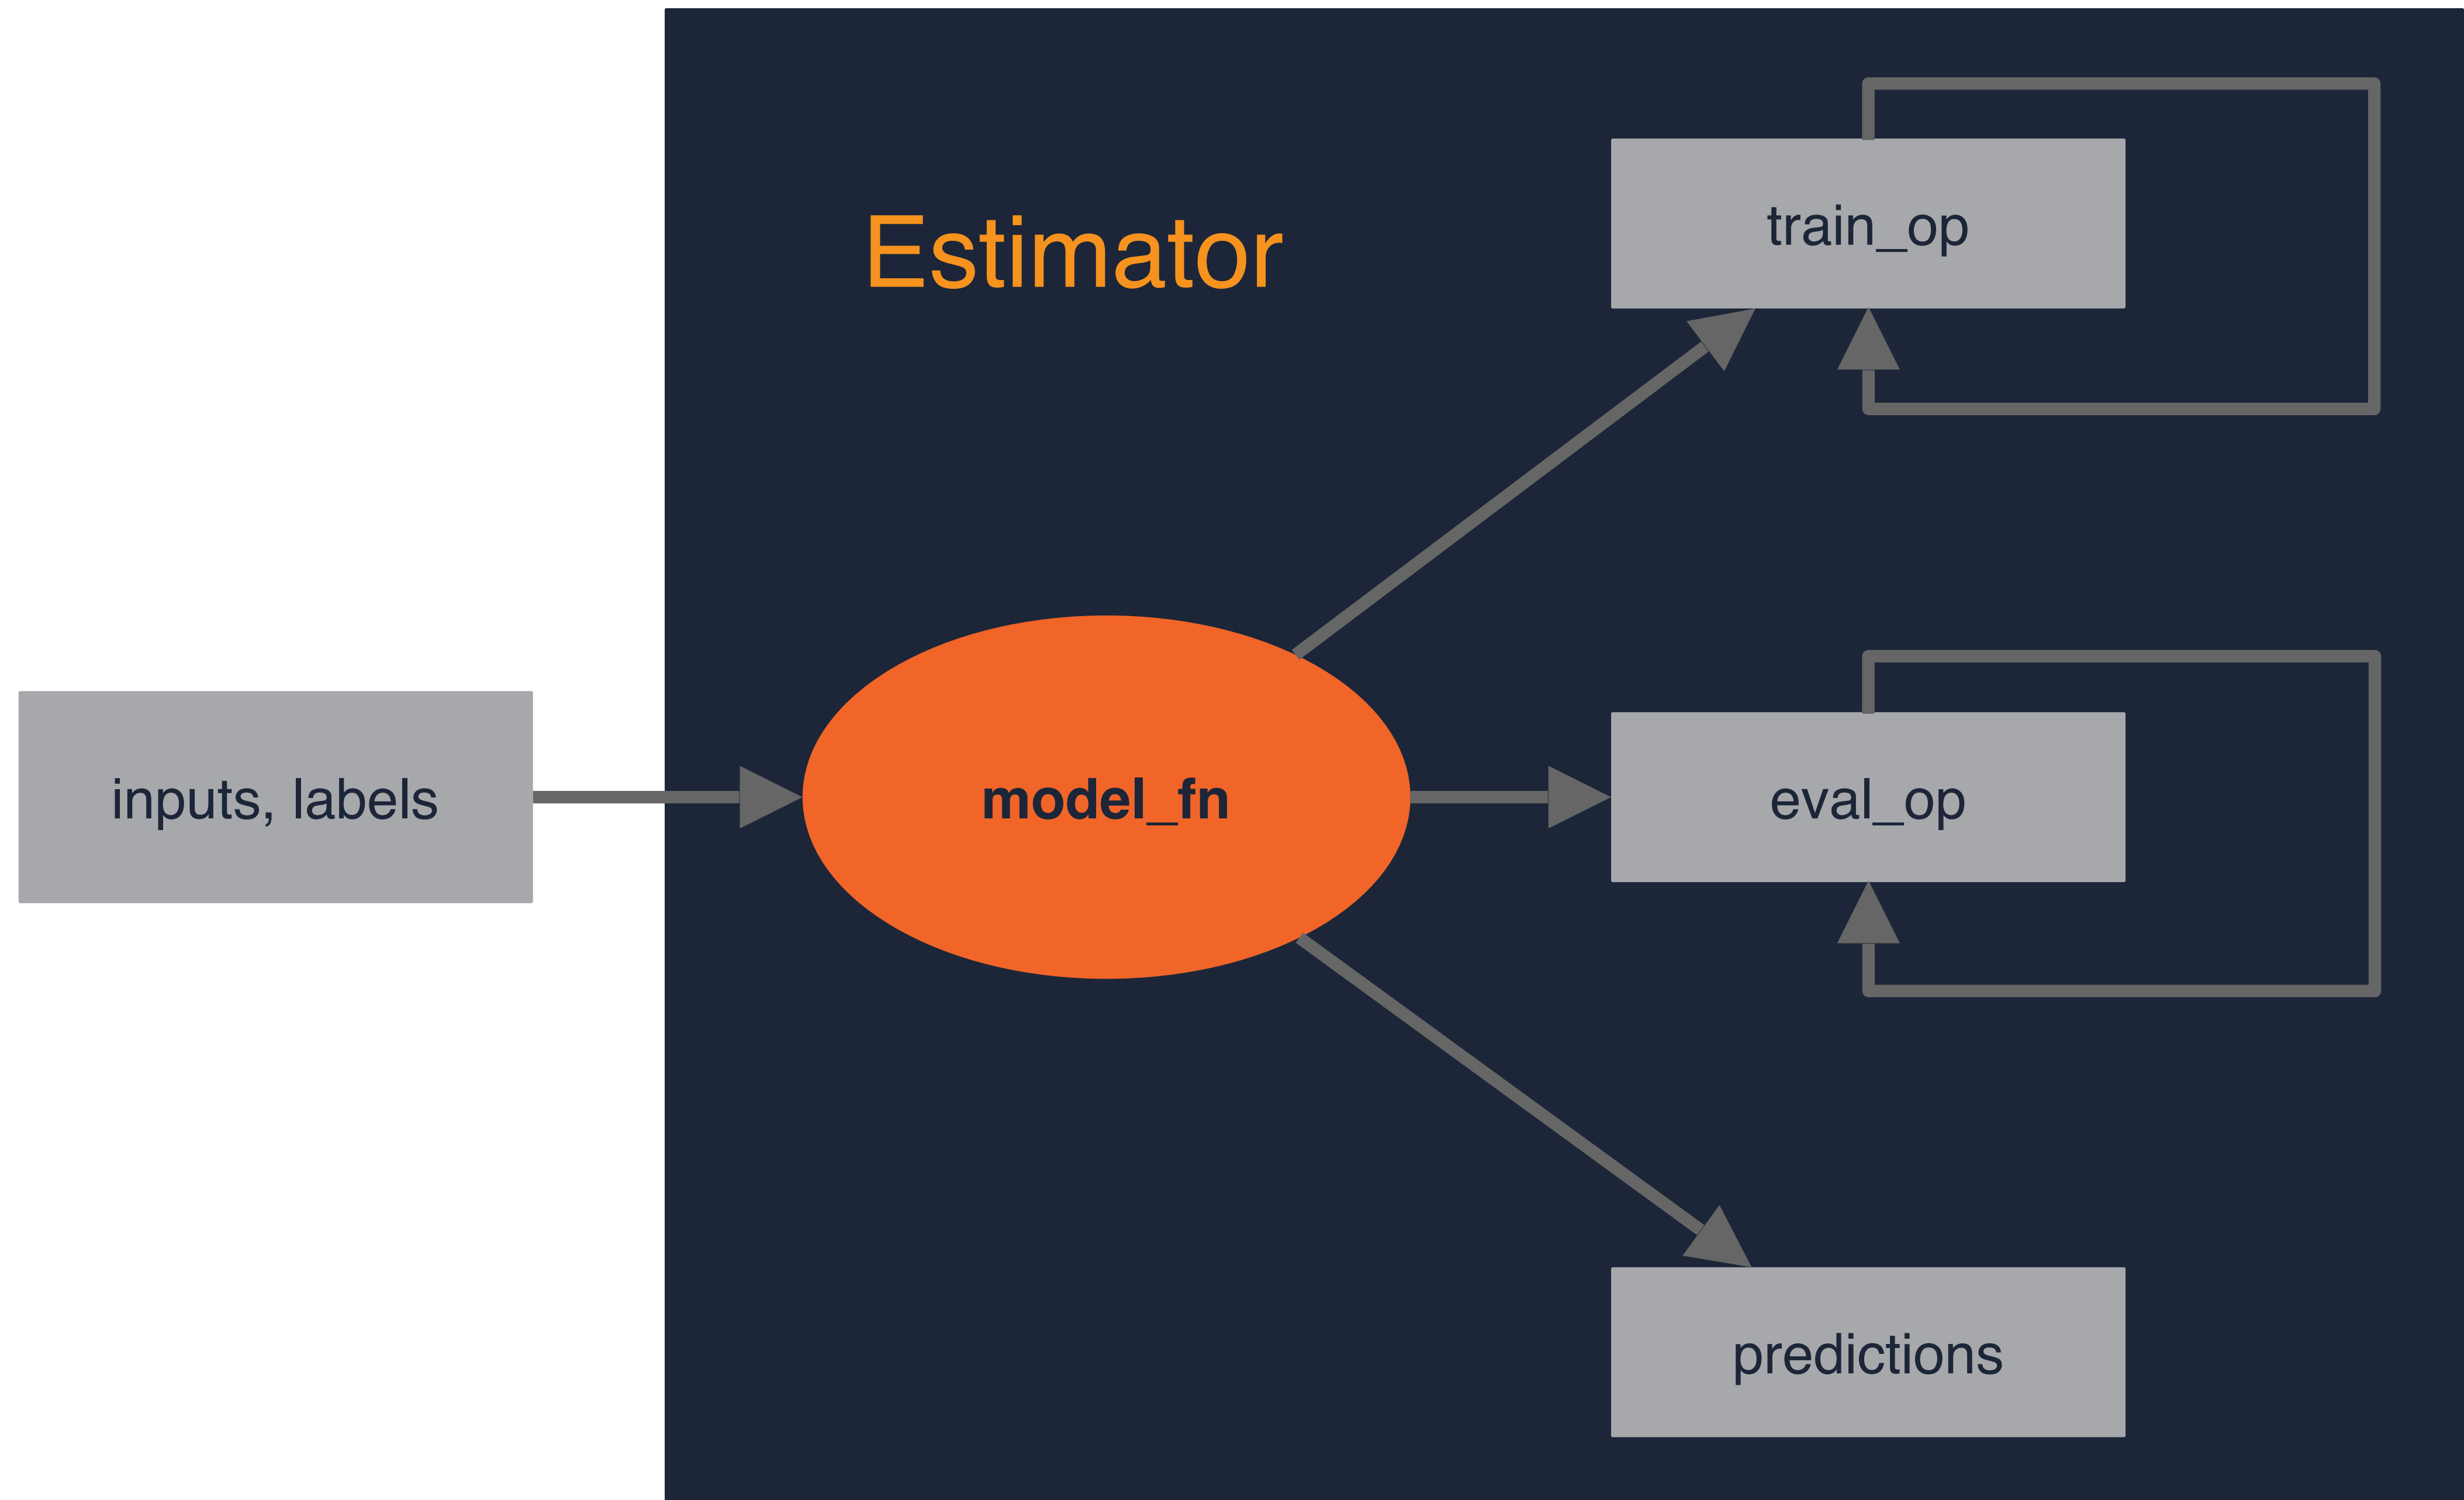
TF Estimators Parts

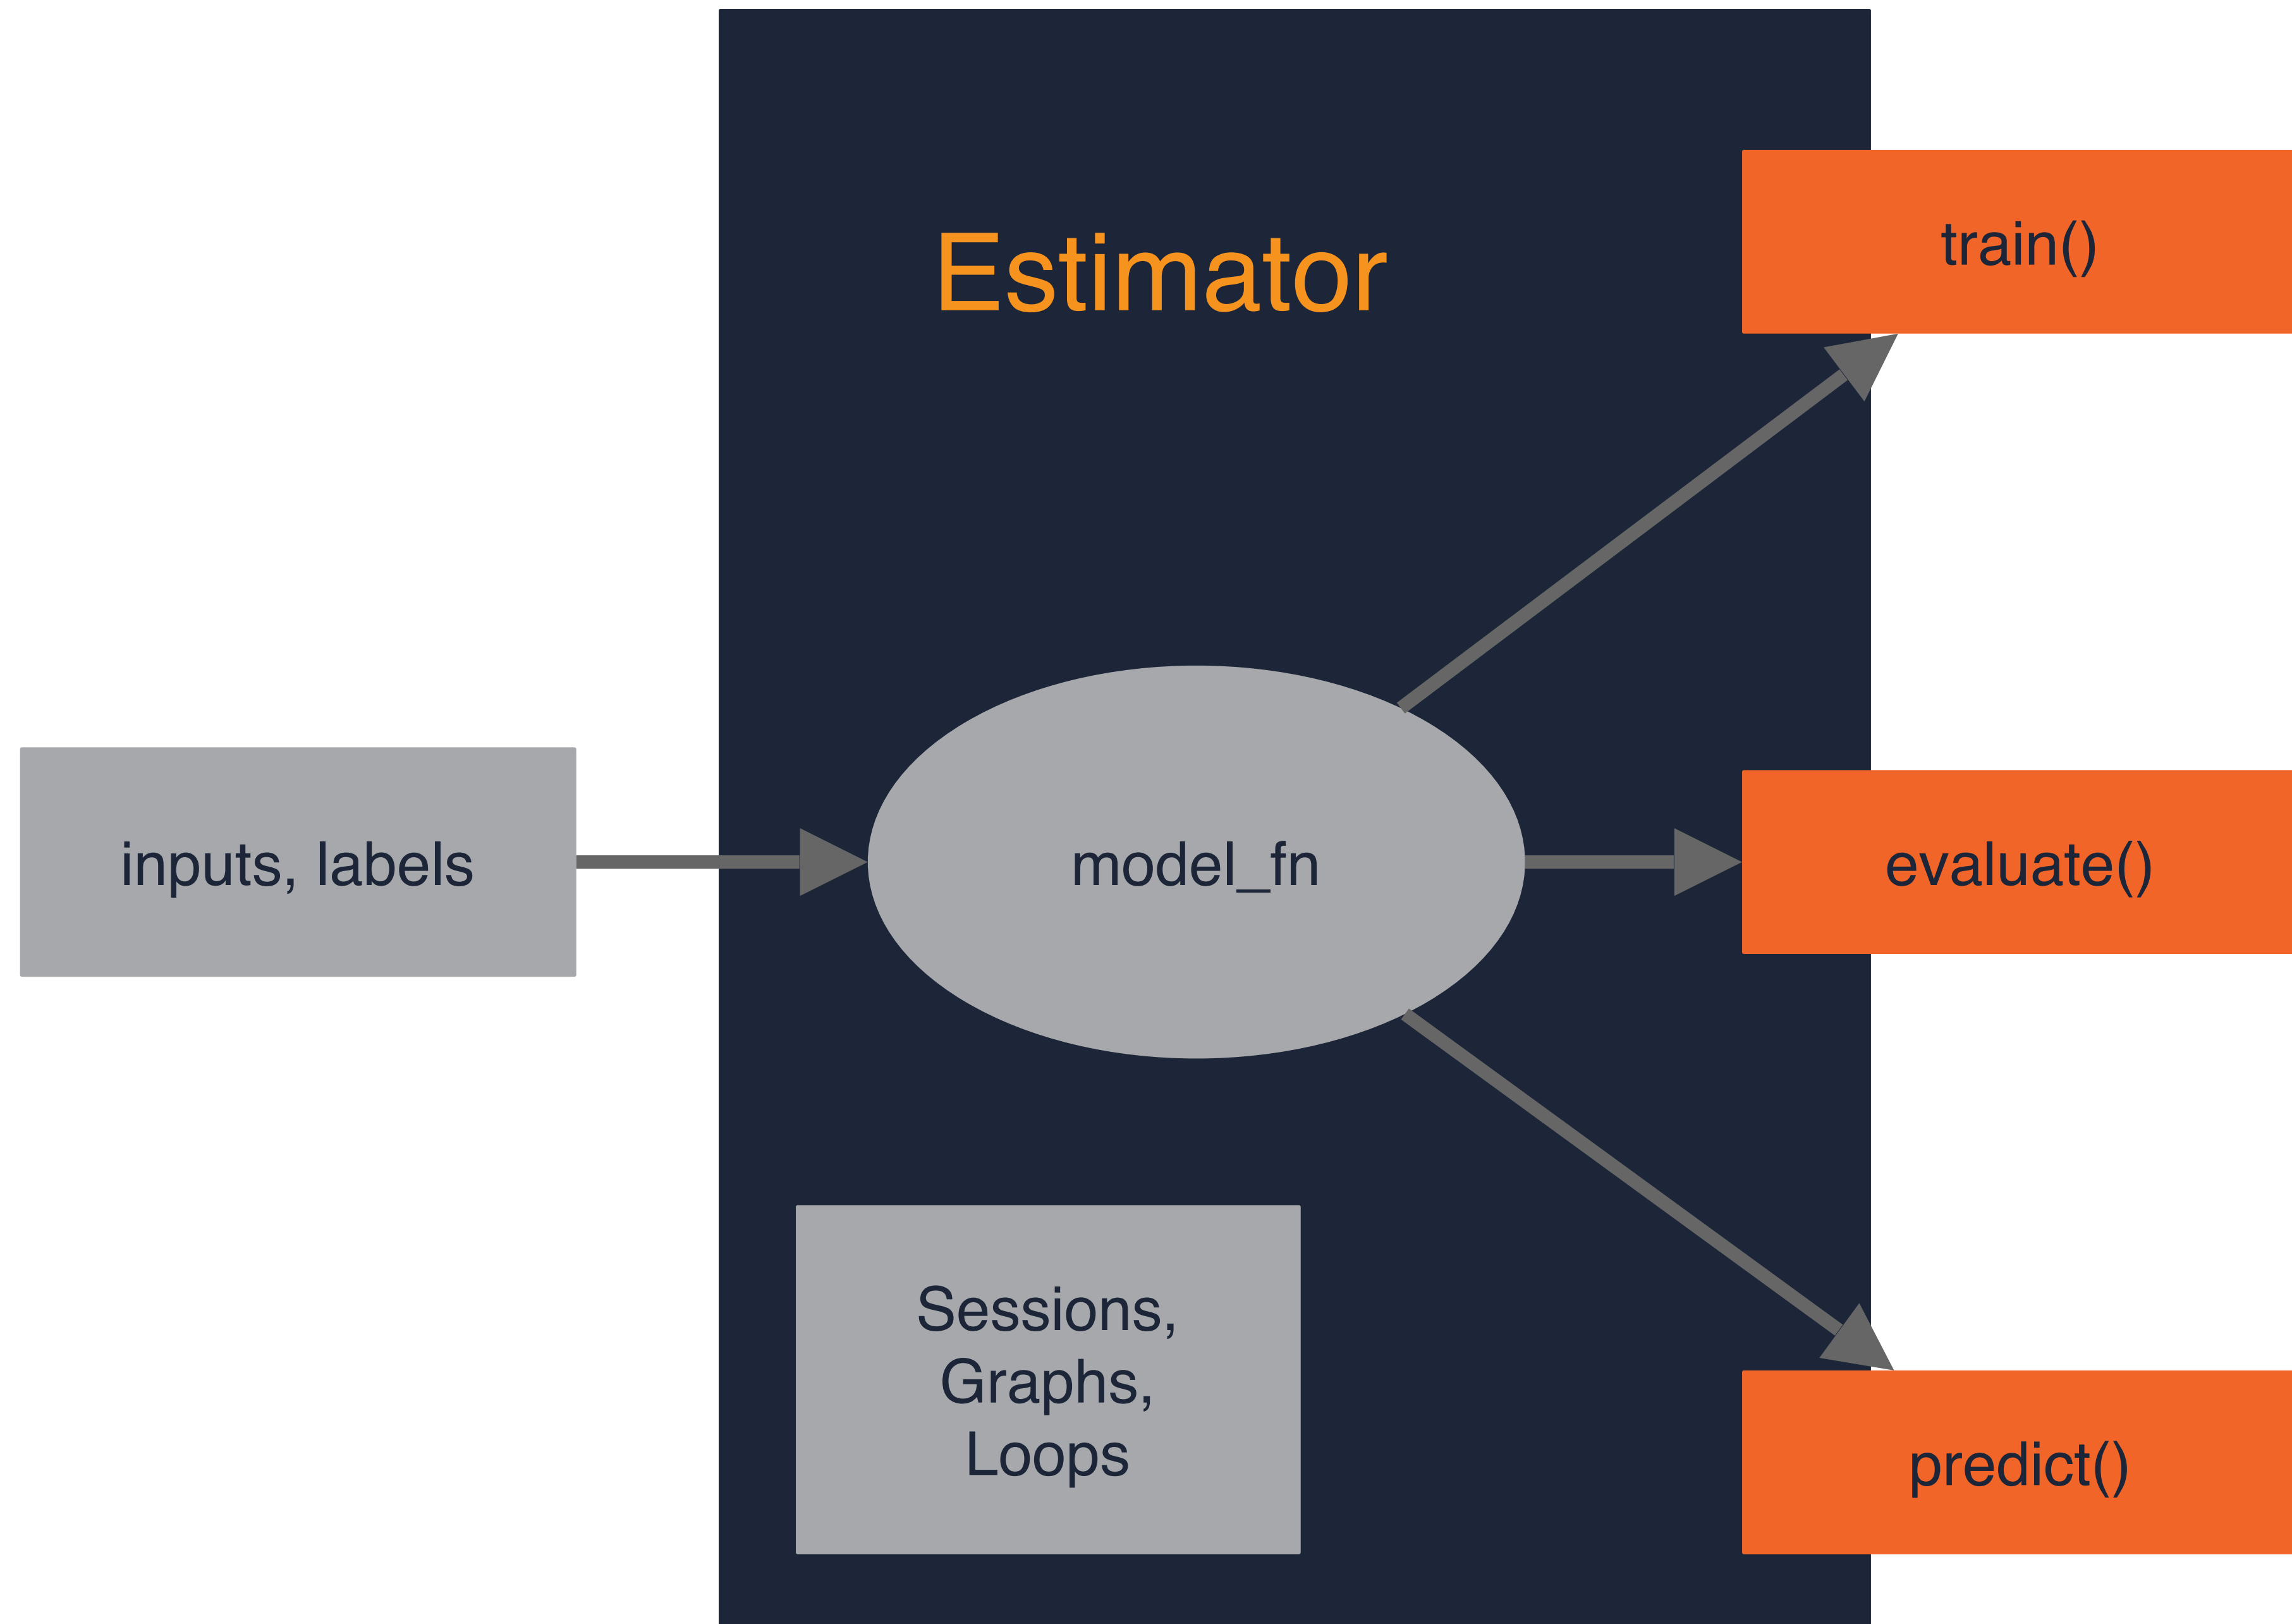
- `input_fn`
- `network_fn`
- `Model_fn`
- Layers
- Datasets
- The functions an Estimator runs

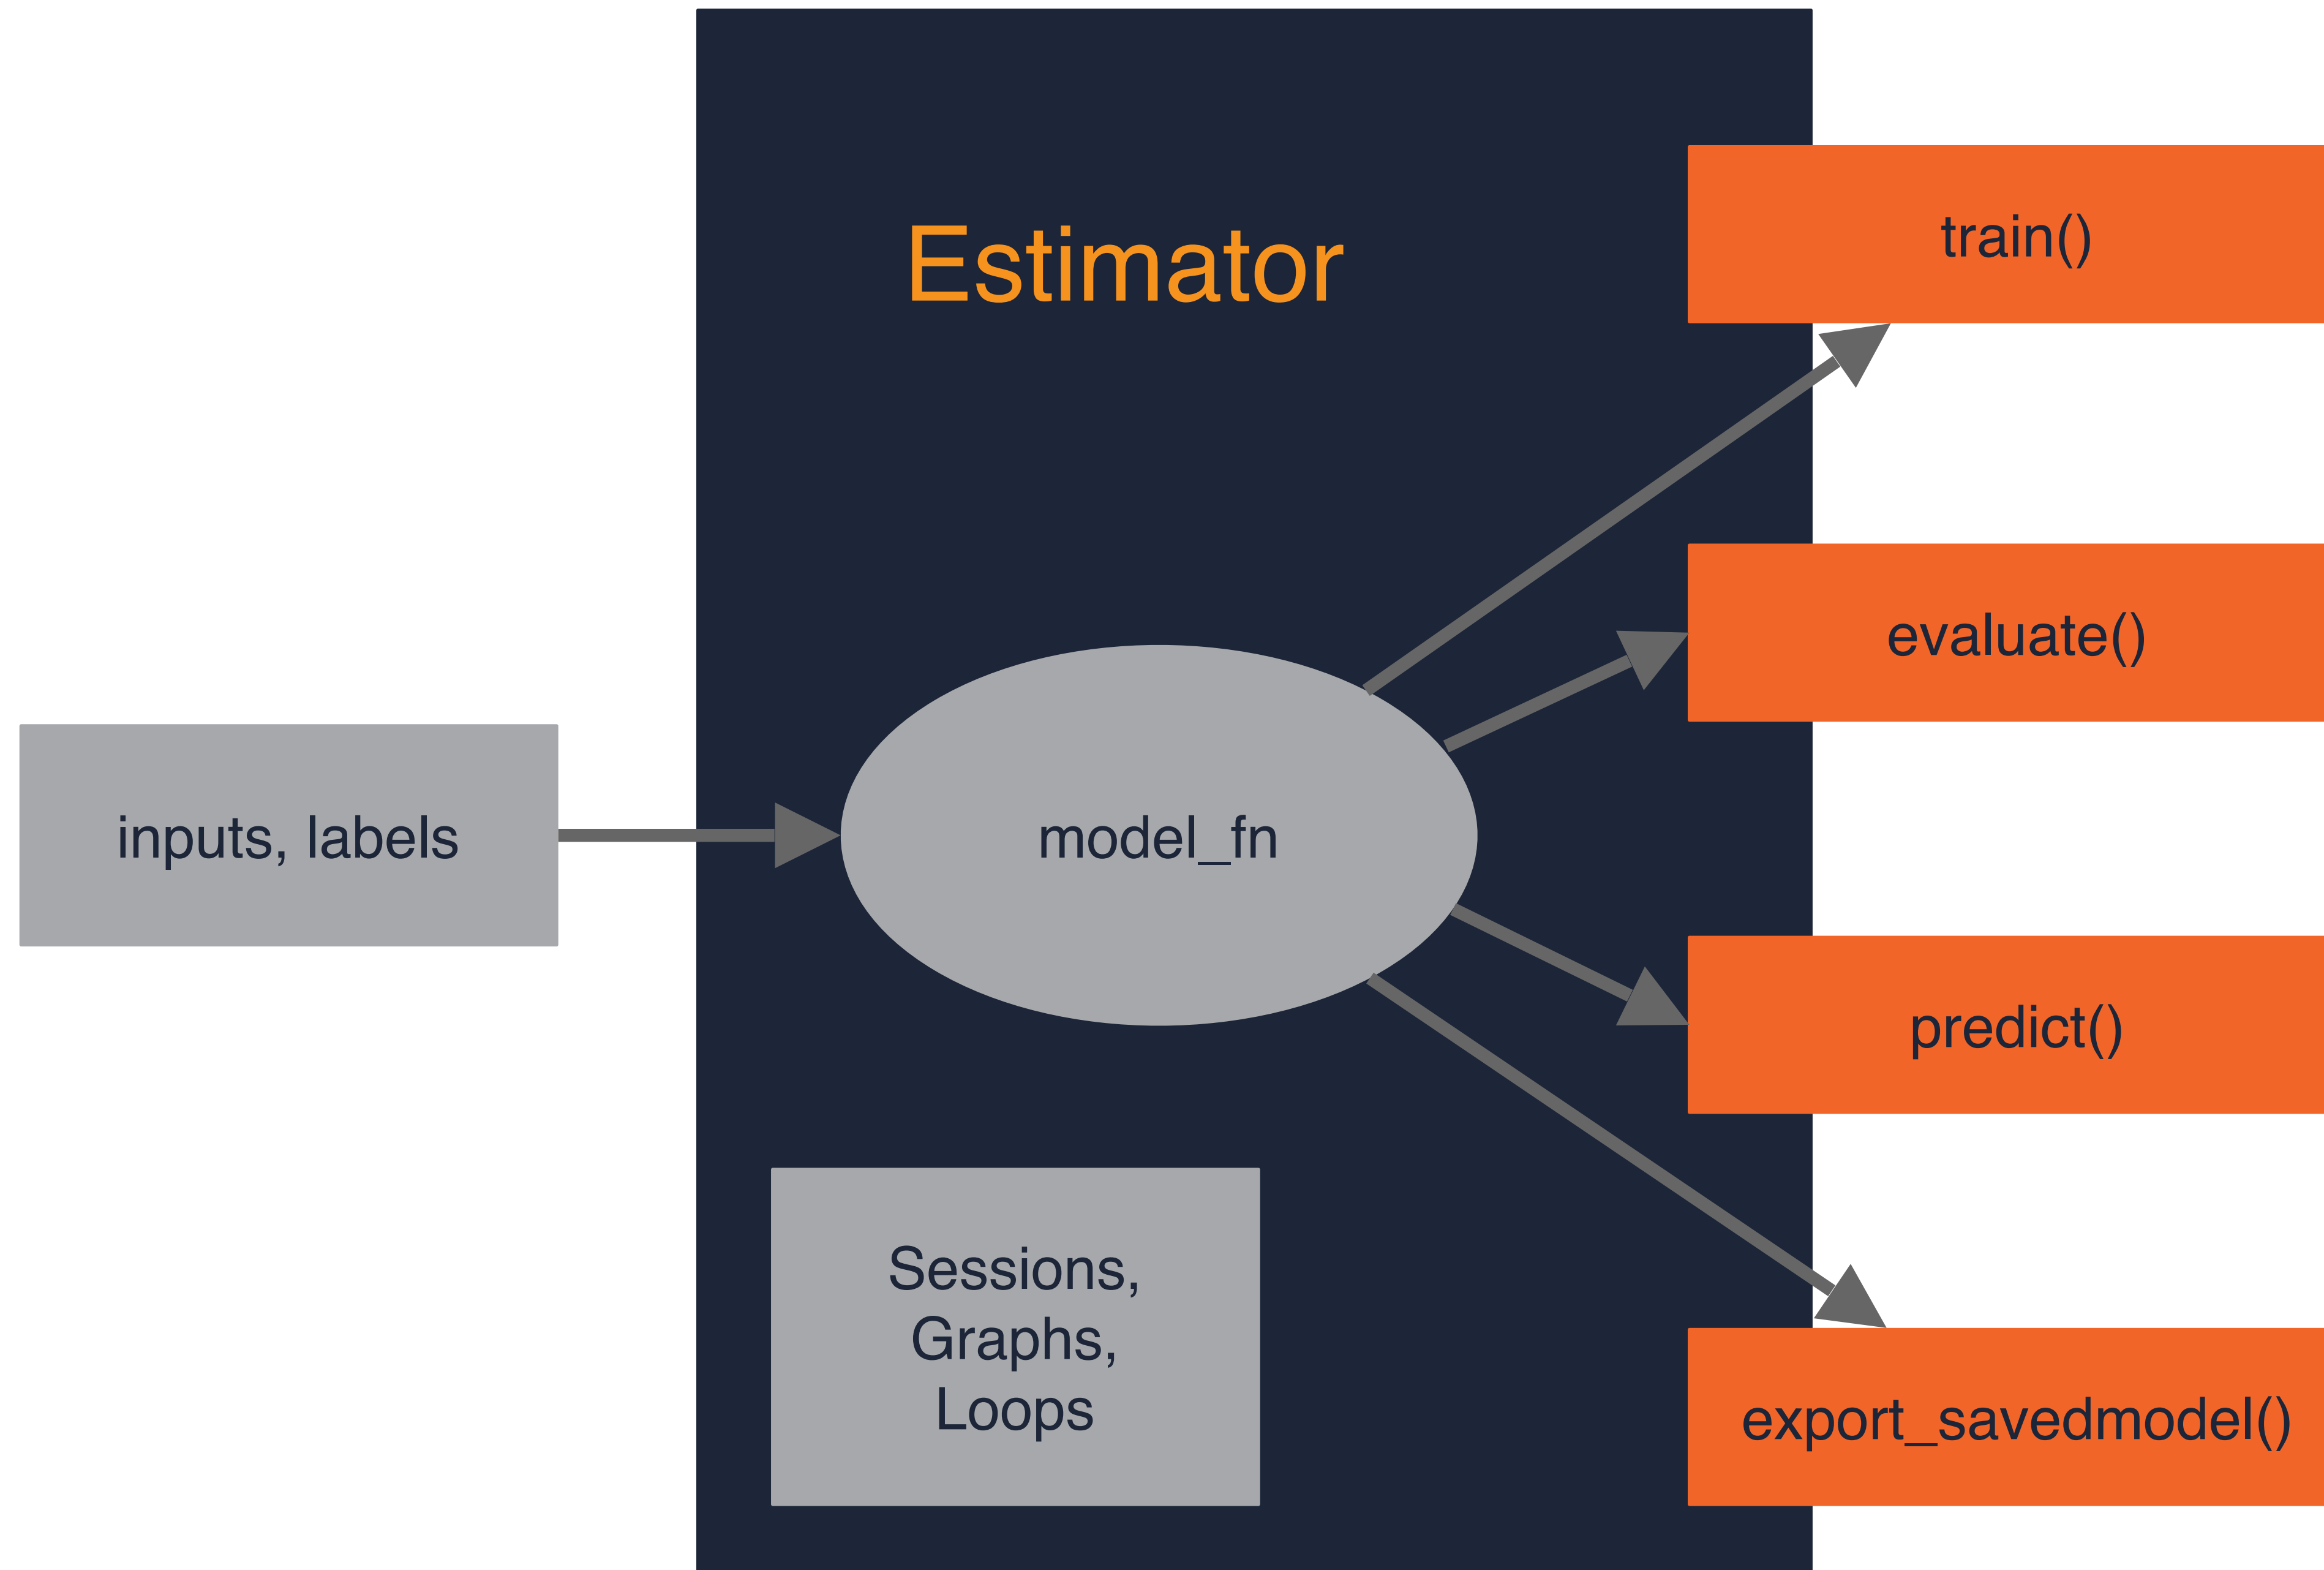




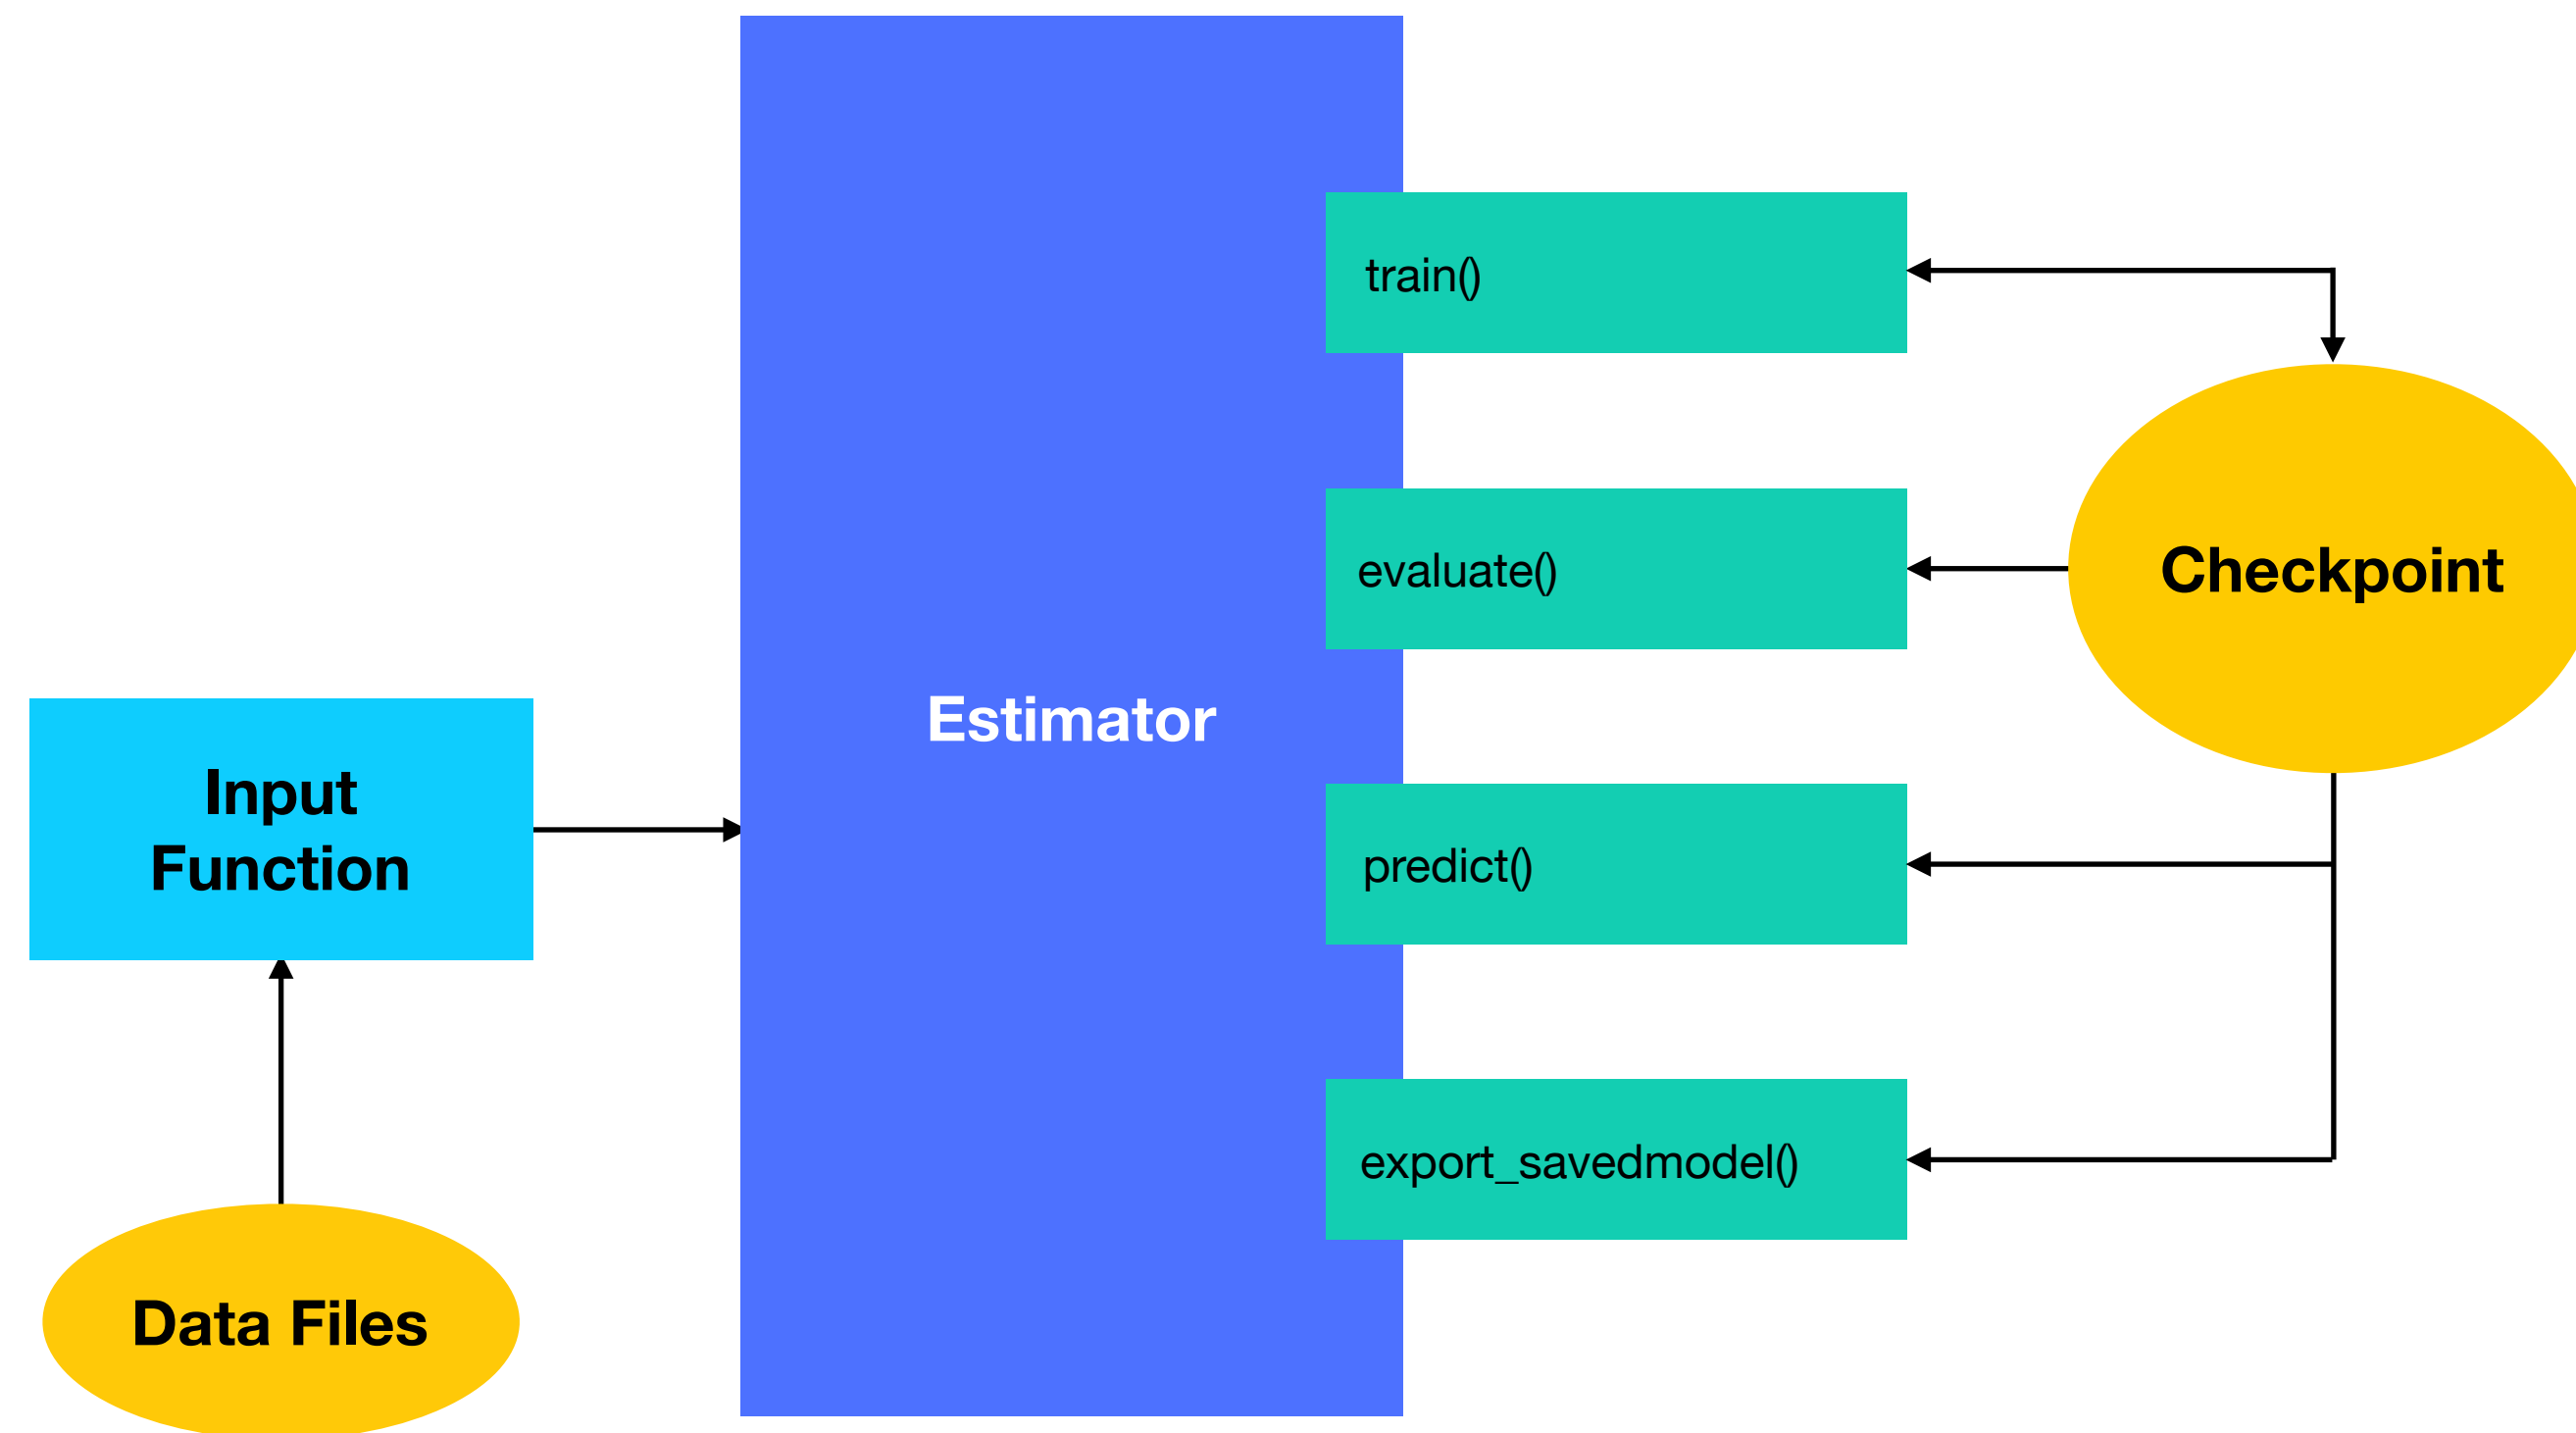








Estimator Framework



Why Estimators ?

- To make it more simple and accessible.
- Allows for better parallelization when training using multi GPUs
- TensorFlow is becoming optimised for it
- It removes the python for loop and converts to C++
- TPUs will only work with Estimators and Datasets APIs

Canned Estimators

- Pre-made Types of networks
- Simple for swapping models and running experiments on the same data
- Currently the types of models are limited
- Probably a lot more in the future
- Possibly trying to take on SKLearn too

Steps

- Load you datasets
- Define column features
- Define your estimator
- .train your estimator
- .evaluate
- .predict
- .export_savedmodel

TensorFlow v1.4

tf.data

tf.keras

custom estimators

Coming soon!!

